

Data Wrangling - Missing Values

Tom Thorpe

March 27, 2018

Springboard Data Science Foundation Class

Exercise 02 - Data Wrangling - Fill in Missing values.

Objective

Practice using *tidy* and *dplyr* packages to clean up data in a practice dataset.

Exercise Results

Load the tidy and dplyr libraries for exercise.

```
#library(devtools)
library(tidy)
library(dplyr)
# install.packages("gdata") needed for read.xls function
# library(gdata) - instructions were to convert to csv before loading, so not needed.
```

Identify input and output files

The practice data set was downloaded into a file called “titanic3.xls”, resaved as “titanic_original.csv” and the cleaned data will be stored in a file named, “titanic_clean.csv”.

```
infile = "C:/Users/Tom/git/datasciencefoundation/DataWrangleExer02/titanic_original.csv"
outfile = "C:/Users/Tom/git/datasciencefoundation/DataWrangleExer02/titanic_clean.csv"
```

Read CSV file into a local dataframe

```
# titanic <- read.xls(infile) %>% tbl_df # instructions were to convert to csv before loading
#titanic <- read.csv(file=infile, header=TRUE, sep=",") %>% tbl_df()
titanic1 <- read.csv(file=infile, header=TRUE, sep=",")
titanic <- tbl_df(titanic1)
```

Lets see what the data looks like:

```
str(titanic)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':   1309 obs. of  14 variables:
## $ pclass : int  1 1 1 1 1 1 1 1 1 1 ...
## $ survived : int  1 1 0 0 0 1 1 0 1 0 ...
## $ name : Factor w/ 1307 levels "Abbing, Mr. Anthony",...: 22 24 25 26 27 31 46 47 51 55 ...
## $ sex : Factor w/ 2 levels "female","male": 1 2 1 2 1 2 1 2 1 2 ...
## $ age : num  29 0.917 2 30 25 ...
## $ sibsp : int  0 1 1 1 1 0 1 0 2 0 ...
```

```
## $ parch      : int   0 2 2 2 2 0 0 0 0 0 ...
## $ ticket     : Factor w/ 929 levels "110152","110413",...: 188 50 50 50 50 125 93 16 77 826 ...
## $ fare       : num   211 152 152 152 152 ...
## $ cabin      : Factor w/ 187 levels "", "A10", "A11",...: 45 81 81 81 81 151 147 17 63 1 ...
## $ embarked   : Factor w/ 4 levels "", "C", "Q", "S": 4 4 4 4 4 4 4 4 2 ...
## $ boat       : Factor w/ 28 levels "", "1", "10", "11",...: 13 4 1 1 1 14 3 1 28 1 ...
## $ body       : int   NA NA NA 135 NA NA NA NA NA 22 ...
## $ home.dest: Factor w/ 370 levels "", "?Havana, Cuba",...: 310 232 232 232 232 238 163 25 23 230 ...
```

```
titanic
```

```
## # A tibble: 1,309 x 14
##   pclass survived name      sex      age sibsp parch ticket  fare cabin
##   <int>    <int> <fct>    <fct>   <dbl> <int> <int> <fct>   <dbl> <fct>
## 1      1      1      1 Allen, Mis~ fema~ 29.0     0     0 24160  211.  B5
## 2      1      1      1 Allison, M~ male  0.917     1     2 113781 152.  C22 ~
## 3      1      0      1 Allison, M~ fema~ 2.00     1     2 113781 152.  C22 ~
## 4      1      0      1 Allison, M~ male 30.0     1     2 113781 152.  C22 ~
## 5      1      0      1 Allison, M~ fema~ 25.0     1     2 113781 152.  C22 ~
## 6      1      1      1 Anderson, ~ male 48.0     0     0 19952  26.6 E12
## 7      1      1      1 Andrews, M~ fema~ 63.0     1     0 13502  78.0 D7
## 8      1      0      1 Andrews, M~ male 39.0     0     0 112050   0.  A36
## 9      1      1      1 Appleton, ~ fema~ 53.0     2     0 11769  51.5 C101
## 10     1      0      1 Artagaveyt~ male 71.0     0     0 PC 17~ 49.5 ""
## # ... with 1,299 more rows, and 4 more variables: embarked <fct>,
## #   boat <fct>, body <int>, home.dest <fct>
```

I noticed that there are three additional variables in the data set than in the data description pointed to by the exercise URL, *boat*, *body* and *home.dest*. I found the full description of the titanic3.xls data set here <http://biostat.mc.vanderbilt.edu/wiki/pub/Main/DataSets/titanic3info.txt>. The *boat* is the lifeboat, *body* is the body identification number, and *home.dest* is the home destination of the passenger.

Default Missing *embarked* data

First fill in any missing data in the *embarked* column with 'S' to represent Southampton. First, Check to see how many are blank.

```
table(titanic$embarked)
```

```
##
##      C    Q    S
##    2 270 123 914
```

There are two empty *embarked* values. Make the updates.

```
titanic$embarked[titanic$embarked==" "|titanic$embarked==" "] <- "S"
#titanic$embarked[grepl("^ !| {0}",titanic$embarked)] <- "S" # couldn't figure out a regular expression
```

Check the results of the change.

```
table(titanic$embarked)
```

```
##
##      C    Q    S
##    0 270 123 916
```

Populate missing *age* data

Populate the missing data with the mean of the other rows with age data. Look at the age data before populating.

```
table(titanic$age)
```

```
##
## 0.1667 0.3333 0.4167 0.6667 0.75 0.8333 0.9167 1 2 3
## 1 1 1 1 3 3 2 10 12 7
## 4 5 6 7 8 9 10 11 11.5 12
## 10 5 6 4 6 10 4 4 1 3
## 13 14 14.5 15 16 17 18 18.5 19 20
## 5 8 2 6 19 20 39 3 29 23
## 20.5 21 22 22.5 23 23.5 24 24.5 25 26
## 1 41 43 1 26 1 47 1 34 30
## 26.5 27 28 28.5 29 30 30.5 31 32 32.5
## 1 30 32 3 30 40 2 23 24 4
## 33 34 34.5 35 36 36.5 37 38 38.5 39
## 21 16 2 23 31 2 9 14 1 20
## 40 40.5 41 42 43 44 45 45.5 46 47
## 18 3 11 18 9 10 21 2 6 14
## 48 49 50 51 52 53 54 55 55.5 56
## 14 9 15 8 6 4 10 8 1 4
## 57 58 59 60 60.5 61 62 63 64 65
## 5 6 3 7 1 5 5 4 5 3
## 66 67 70 70.5 71 74 76 80
## 1 1 2 1 2 1 1 1
```

```
summary(titanic$age)
```

```
## Min. 1st Qu. Median Mean 3rd Qu. Max. NA's
## 0.1667 21.0000 28.0000 29.8811 39.0000 80.0000 263
```

```
filter(titanic,is.na(age)) %>% count
```

```
## # A tibble: 1 x 1
## n
## <int>
## 1 263
```

```
select(titanic,name,age) %>% print(n=18)
```

```
## # A tibble: 1,309 x 2
## name age
## <fct> <dbl>
## 1 Allen, Miss. Elisabeth Walton 29.0
## 2 Allison, Master. Hudson Trevor 0.917
## 3 Allison, Miss. Helen Loraine 2.00
## 4 Allison, Mr. Hudson Joshua Creighton 30.0
## 5 Allison, Mrs. Hudson J C (Bessie Waldo Daniels) 25.0
## 6 Anderson, Mr. Harry 48.0
## 7 Andrews, Miss. Kornelia Theodosia 63.0
## 8 Andrews, Mr. Thomas Jr 39.0
## 9 Appleton, Mrs. Edward Dale (Charlotte Lamson) 53.0
## 10 Artagaveytia, Mr. Ramon 71.0
## 11 Astor, Col. John Jacob 47.0
```

```
## 12 Astor, Mrs. John Jacob (Madeleine Talmadge Force) 18.0
## 13 Aubart, Mme. Leontine Pauline 24.0
## 14 "Barber, Miss. Ellen \"Nellie\"" 26.0
## 15 Barkworth, Mr. Algernon Henry Wilson 80.0
## 16 Baumann, Mr. John D NA
## 17 Baxter, Mr. Quigg Edmond 24.0
## 18 Baxter, Mrs. James (Helene DeLaudeniére Chaput) 50.0
## # ... with 1,291 more rows
```

Notice that the 16th entry has NA for age. Set the empty age data to the mean of the non-empty age.

```
meanAge= mean(titanic$age,na.rm=TRUE)
medianAge= median(titanic$age,na.rm=TRUE)
# meanAge= mean(titanic$age[titanic$age>=0],na.rm=TRUE)
meanAge
```

```
## [1] 29.88113
```

```
medianAge
```

```
## [1] 28
```

```
meanAge<-trunc(meanAge) # truncate the age since only age less than 1 have fractional values.
#titanic$age[titanic$age==NA] <- meanAge
titanic$age[is.na(titanic$age)] <- meanAge

#Check the results
select(titanic,name,age) %>% print(n=18)
```

```
## # A tibble: 1,309 x 2
##   name age
##   <fct> <dbl>
## 1 Allen, Miss. Elisabeth Walton 29.0
## 2 Allison, Master. Hudson Trevor 0.917
## 3 Allison, Miss. Helen Loraine 2.00
## 4 Allison, Mr. Hudson Joshua Creighton 30.0
## 5 Allison, Mrs. Hudson J C (Bessie Waldo Daniels) 25.0
## 6 Anderson, Mr. Harry 48.0
## 7 Andrews, Miss. Kornelia Theodosia 63.0
## 8 Andrews, Mr. Thomas Jr 39.0
## 9 Appleton, Mrs. Edward Dale (Charlotte Lamson) 53.0
## 10 Artagaveytia, Mr. Ramon 71.0
## 11 Astor, Col. John Jacob 47.0
## 12 Astor, Mrs. John Jacob (Madeleine Talmadge Force) 18.0
## 13 Aubart, Mme. Leontine Pauline 24.0
## 14 "Barber, Miss. Ellen \"Nellie\"" 26.0
## 15 Barkworth, Mr. Algernon Henry Wilson 80.0
## 16 Baumann, Mr. John D 29.0
## 17 Baxter, Mr. Quigg Edmond 24.0
## 18 Baxter, Mrs. James (Helene DeLaudeniére Chaput) 50.0
## # ... with 1,291 more rows
```

```
filter(titanic,is.na(age)) %>% count
```

```
## # A tibble: 1 x 1
##       n
##   <int>
```

```
## 1      0
```

How else could the missing age data be populated? The median is an alternative measure that may be closer to a representative age. The median in this case is 28. The mean is 29.88 rounded to 30. I believe that since the median is less than the mean, there are more younger people than old people on the ship. I would tend to use the median over the mean in this case. Somehow it just feels better. Maybe I will have a better answer after taking the statistics part of the course.

Populate missing Lifeboat data

There is missing data for the lifeboat column.

Does it make sense to fill in the cabin numbers with a value? To me, it does not make sense to try to come up with an estimate of the cabin number. How could any meaningful data be created? It makes sense to set the missing data to “NONE”.

What does a missing value mean here? Missing data could mean the records for the cabin assignment were lost or the people were assigned to a big dormitory area instead of a room. But there was probably more than one dormitory and the specific dormitory is unknown. Either way, the best that can be done is to do as the assignment suggests and assign missing data the value of “NONE”.

Update the missing data in the lifeboat data with “NONE”. I tried using these commands

```
titanic$boat[titanic$boat==""] <- as.factor("NONE")
```

```
titanic$boat[titanic$boat==""] <- "NONE"
```

But both resulted in the following warning:

Warning in [`<-`].factor(*tmp*, titanic\$boat == "", value = structure(c(13L, : invalid factor level, NA generated

So I did some research and converted the *boat* column from a factor to character data. But I don’t understand why the column was a factor and not character data to begin with. I need to research this.

```
filter(titanic, boat == "") %>% count()
```

```
## # A tibble: 1 x 1
##       n
##   <int>
## 1    823
```

```
titanic$boat <- as.character.factor(titanic$boat)
titanic$boat[titanic$boat==""] <- "NONE"
#Check the results
select(titanic,name,boat) %>% print(n=15)
```

```
## # A tibble: 1,309 x 2
##       name                                boat
##   <fct>                                <chr>
## 1 Allen, Miss. Elisabeth Walton          2
## 2 Allison, Master. Hudson Trevor        11
## 3 Allison, Miss. Helen Loraine         NONE
## 4 Allison, Mr. Hudson Joshua Creighton  NONE
## 5 Allison, Mrs. Hudson J C (Bessie Waldo Daniels) NONE
## 6 Anderson, Mr. Harry                   3
## 7 Andrews, Miss. Kornelia Theodosia    10
## 8 Andrews, Mr. Thomas Jr               NONE
## 9 Appleton, Mrs. Edward Dale (Charlotte Lamson) D
```

```
## 10 Artagaveytia, Mr. Ramon      NONE
## 11 Astor, Col. John Jacob      NONE
## 12 Astor, Mrs. John Jacob (Madeleine Talmadge Force) 4
## 13 Aubart, Mme. Leontine Pauline 9
## 14 "Barber, Miss. Ellen \"Nellie\" 6
## 15 Barkworth, Mr. Algernon Henry Wilson B
## # ... with 1,294 more rows
```

```
filter(titanic, boat == "") %>% count()
```

```
## # A tibble: 1 x 1
##       n
##   <int>
## 1     0
```

Create has_cabin_number column

Create a *has_cabin_number* column with a value of 1 if the *cabin* column is non-blank.

```
titanic <- mutate(titanic, has_cabin_number = ifelse(cabin != "", 1, 0))
# Check results
select(titanic, name, cabin, has_cabin_number) %>% print(n=15)
```

```
## # A tibble: 1,309 x 3
##   name                cabin has_cabin_number
##   <fct>              <fct>         <dbl>
## 1 Allen, Miss. Elisabeth Walton      B5             1.
## 2 Allison, Master. Hudson Trevor    C22 C~           1.
## 3 Allison, Miss. Helen Loraine      C22 C~           1.
## 4 Allison, Mr. Hudson Joshua Creighton C22 C~           1.
## 5 Allison, Mrs. Hudson J C (Bessie Waldo Daniels) C22 C~           1.
## 6 Anderson, Mr. Harry                E12             1.
## 7 Andrews, Miss. Kornelia Theodosia D7              1.
## 8 Andrews, Mr. Thomas Jr            A36             1.
## 9 Appleton, Mrs. Edward Dale (Charlotte Lamson) C101            1.
## 10 Artagaveytia, Mr. Ramon          ""              0.
## 11 Astor, Col. John Jacob            C62 C~           1.
## 12 Astor, Mrs. John Jacob (Madeleine Talmadge For~ C62 C~           1.
## 13 Aubart, Mme. Leontine Pauline      B35             1.
## 14 "Barber, Miss. Ellen \"Nellie\"    ""              0.
## 15 Barkworth, Mr. Algernon Henry Wilson A23             1.
## # ... with 1,294 more rows
```

```
glimpse(titanic)
```

```
## Observations: 1,309
## Variables: 15
## $ pclass      <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
## $ survived    <int> 1, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 1, 1, 1, ...
## $ name        <fct> Allen, Miss. Elisabeth Walton, Allison, Maste...
## $ sex         <fct> female, male, female, male, female, male, fem...
## $ age         <dbl> 29.0000, 0.9167, 2.0000, 30.0000, 25.0000, 48...
## $ sibsp       <int> 0, 1, 1, 1, 1, 0, 1, 0, 2, 0, 1, 1, 0, 0, ...
## $ parch       <int> 0, 2, 2, 2, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
## $ ticket      <fct> 24160, 113781, 113781, 113781, 113781, 19952,...
## $ fare        <dbl> 211.3375, 151.5500, 151.5500, 151.5500, 151.5...
```

```
## $ cabin          <fct> B5, C22 C26, C22 C26, C22 C26, C22 C26, E12, ...
## $ embarked       <fct> S, S, S, S, S, S, S, S, S, C, C, C, C, S, S, ...
## $ boat           <chr> "2", "11", "NONE", "NONE", "NONE", "3", "10",...
## $ body           <int> NA, NA, NA, 135, NA, NA, NA, NA, NA, 22, 124,...
## $ home.dest       <fct> St Louis, MO, Montreal, PQ / Chesterville, ON...
## $ has_cabin_number <dbl> 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, ...
```

Save the cleaned data

```
write.csv(titanic, file=outfile,row.names=FALSE)
```

That concludes the exercise.